

Kimberly Huang

New York, NY | 332-999-4451 | huang0818hqj@outlook.com | [linkedin.com/in/qijinghuang](https://www.linkedin.com/in/qijinghuang) | kimberly-huang.vercel.app

Applied AI and data engineer finishing an M.S. at Columbia (Dec 2026). I ship AI systems for operational work — multimodal extraction pipelines, multi-agent workflows, document automation — pairing model output with deterministic validation and an explicit human review boundary.

EDUCATION

Columbia University — M.S. in Applied Analytics <i>GPA 4.00 Applied Generative AI, Machine Learning, Data Engineering, Database Systems</i>	Sep 2025 – Dec 2026 <i>New York, NY</i>
University of Dundee — B.S. in Applied Finance <i>Top 1% of cohort 2022–2024 Global Excellence Scholarship ACCA certified</i>	Sep 2020 – Jun 2024 <i>Dundee, United Kingdom</i>

TECHNICAL SKILLS

AI & Agents: LLM APIs (Claude, OpenAI), multi-agent orchestration, agentic workflows, RAG, LangChain, human-in-the-loop design
Programming: Python (pandas, NumPy, Pillow, pypdf, openpyxl), SQL, TypeScript, JavaScript, Node.js, R, Git, GitHub Actions CI
Data & Cloud: PostgreSQL, MongoDB, BigQuery, ETL and pipeline design, AWS, Vercel, GraphQL, OAuth2, MQTT, PPTX/DOCX/OOXML automation
Analytics: Power BI, Tableau, Looker, Excel, A/B testing, logistic regression, EDA, data modeling | **Languages:** English, Mandarin Chinese

EXPERIENCE

Deutsche Windtechnik <i>AI & Data Engineer Intern</i>	Jun 2026 – Sep 2026 <i>Houston, TX</i>
- Shipped the first automated detection capability for blade-repair photo reuse — an operational gap no existing tooling covered — via a Python image-hashing pipeline that surfaced cross-tower reuse risk to operations - Processed 171 supplier PDFs and 7,248 embedded image objects with zero read errors, screening 4,480 candidate photos by SHA-256 hash into 141 exact-duplicate groups across three risk tiers, with structured Excel output for review - Established that SharePoint knowledge sources index embedded PDFs as text only, then re-architected an infeasible agent-only design into a hybrid workflow: local Python for image processing, internal LLM agent for QA summarization - Packaged the pipeline for non-technical Windows users with a batch runner, dependency auto-install and OneDrive-aware file-lock fallbacks; reported methodology and flagged findings to US and Germany stakeholders	
Temu (PDD Holdings) <i>Data Analyst, Growth & Operations</i>	Jun 2024 – Jul 2025 <i>Shenzhen, China</i>
- Built SQL extraction pipelines and Power BI dashboards tracking 15+ KPIs across 20+ global regions, cutting manual reporting effort 40% and giving growth and operations one source of truth for GMV and margin variance - Designed data models for purchasing trends and category-level margin performance, feeding quarterly forecasting and translating funnel behavior into pricing and category decisions - Ran A/B experiments on promotional variants that lifted new-user activation 12%, presenting results to senior cross-functional stakeholders while managing 20+ concurrent data workstreams	
Holt Green <i>Data & Research Analyst Intern</i>	Jun 2023 – Aug 2023
Analyzed financial datasets and compiled the supporting evidence into structured research reports for leadership	

PROJECTS

Student Signal Copilot — Multi-Agent Facilities Workflow <i>github.com/ivelf/Student-Signal-Copilot 3rd Place Overall, NYC Tech Week 2026</i>	Jun 2026 <i>New York, NY</i>
- Architected a supervisor-orchestrated pipeline of 7 specialist agents turning one-sentence field reports into confidence-labeled, compliance-checked work orders grounded in verified asset records, with tiered Claude models (Haiku per specialist, Sonnet for cross-signal overview) bounding cost by design - Kept two of the four workflow stages rule-based and zero-LLM and required human execution for every real-world action, labeling conclusions Verified / Likely / Inferred / Missing so model uncertainty stays visible - Integrated the CriticalAsset GraphQL API with NYC DOB violation and 311 records over live SSE streaming, with OAuth2 token caching, auto-refresh and server-side credential isolation so keys never reach the browser	
DWT Document Agent — Internal AI Document Platform <i>github.com/Kimberly-Huang/dwt-document-agent Python · Node.js · OpenAI API · Vercel</i>	Jul 2026 <i>Houston, TX</i>
- Built and deployed an internal platform converting natural-language requests and uploaded Office, PDF and photo sources into editable, brand-compliant PPTX and DOCX across four business workflows — private Vercel Blob storage, signed-cookie pilot auth and 78 automated tests on CI - Designed a hybrid engine: the model writes narrative and per-slide scene composition while deterministic code enforces brand tokens, validates geometry, binds source data to native editable charts and falls back to verified layouts - Resolved production failures across OOXML list-style resolution and node ordering, Vercel's read-only serverless filesystem, Blob SDK version drift and Unicode-safe download streaming	
EventScope — Financial Event Intelligence Platform <i>github.com/Kimberly-Huang/eventscope Python · SQL · PostgreSQL</i>	Oct – Dec 2025 <i>New York, NY</i>
Built a Python and PostgreSQL event-analytics platform over 65,000+ SEC filing data points, correlating earnings surprises with abnormal returns through a schema that keeps ingestion, storage and analysis cleanly separated	